



Varun Sagar
Theegala

SQL OPTIMIZATION

NOT SURE IF MY QUERY IS FAST



OR SLOW

CAN BE HELPFUL!

LEARN HOW →



Varun Sagar
Theegala

Efficient SQL querying is key to streamlined data processing in organizations.

As a result, there are many **simple optimization practices** that have been adopted, **vital for enhancing database performance, reducing resource load**, and delivering quicker insights.



Varun Sagar
Theegala

**LET'S DIVE DEEP INTO 10 SIMPLE
OPTIMIZATION PRACTICES**



#1

INDEXING

WHAT IT IS ?

Using indexes to improve the database query performance.

WHY IT MATTERS ?

Indexes speed up the retrieval of rows from a table, making queries more efficient.

HOW TO OPTIMIZE ?

Identify frequently queried columns and create indexes on those columns.



LET'S WORK WITH AN EXAMPLE

Before Optimization:

sql

```
SELECT first_name, last_name  
FROM employees  
WHERE department_id = 50;
```

After Optimization:

sql

```
-- Assuming an index is created on department_id  
CREATE INDEX idx_department ON employees(department_id);  
  
SELECT first_name, last_name  
FROM employees  
WHERE department_id = 50;
```

**Making department_id an index
column will impact query speed**



Varun Sagar
Theegala

NEXT ➡

#2

Avoiding SELECT *

WHAT IT IS ?

Specifying needed columns instead of using
SELECT *

WHY IT MATTERS ?

Increases query efficiency by reducing data
load and processing time.

HOW TO OPTIMIZE ?

Analyze the data requirements and explicitly
list only the necessary columns in your
SELECT statement.



LET'S WORK WITH AN EXAMPLE

Before Optimization:

sql

```
SELECT *  
FROM employees;
```

After Optimization:

sql

```
SELECT first_name, last_name, email  
FROM employees;
```



#3

Use Joins Instead of Sub-queries

WHAT IT IS ?

Replacing sub-queries with joins to enhance performance.

WHY IT MATTERS ?

Joins are generally more efficient and faster than nested sub-queries.

HOW TO OPTIMIZE ?

Convert sub-queries into JOIN operations where possible to ensure the most efficient type of JOIN for your data.



LET'S WORK WITH AN EXAMPLE

Before Optimization:

sql

```
SELECT e.employee_id, e.first_name
FROM employees e
WHERE e.department_id IN
      (SELECT department_id
       FROM departments
       WHERE manager_id > 200);
```

After Optimization:

sql

```
SELECT e.employee_id, e.first_name
FROM employees e
JOIN departments d
      ON e.department_id = d.department_id
WHERE d.manager_id > 200;
```



#4

Proper Use Of WHERE Clauses

WHAT IT IS ?

Efficiently filtering data using WHERE clauses.

WHY IT MATTERS ?

Reduces the amount of data processed, improving query speed.

HOW TO OPTIMIZE ?

Utilize WHERE clauses to filter data as early as possible in the query process.



LET'S WORK WITH AN EXAMPLE

Before Optimization:

sql

```
SELECT first_name, last_name  
FROM employees  
ORDER BY employee_id  
FETCH FIRST 10 ROWS ONLY;
```

After Optimization:

sql

```
SELECT first_name, last_name  
FROM employees  
WHERE hire_date > '2005-01-01'  
ORDER BY employee_id;
```



#5

Limiting Result Set

WHAT IT IS ?

Using clauses like LIMIT or FETCH to restrict the number of rows returned.

WHY IT MATTERS ?

Prevents over-fetching of data, saving resources and time.

HOW TO OPTIMISE ?

Implement LIMIT or FETCH FIRST clauses in queries where the full dataset is not required.



LET'S WORK WITH AN EXAMPLE

Before Optimization:

sql

```
SELECT * FROM employees;
```

After Optimization:

sql

```
SELECT employee_id, first_name, last_name  
FROM employees  
FETCH FIRST 100 ROWS ONLY;
```

As shown, entering column names to pull would reduce the response time



Varun Sagar
Theegala

NEXT ➡

#6

Optimize GROUP BY & HAVING Clauses

WHAT IT IS ?

Efficiently grouping data and filtering groups.

WHY IT MATTERS ?

Improves performance, especially in large datasets.

HOW TO OPTIMIZE ?

Ensure the use of GROUP BY and HAVING clauses is done in a way that minimizes the amount of data being grouped.



LET'S WORK WITH AN EXAMPLE

Before Optimization:

sql

```
SELECT department_id, COUNT(*)  
FROM employees  
GROUP BY department_id  
HAVING COUNT(*) > 5;
```

After Optimization:

sql

```
SELECT department_id, COUNT(*)  
FROM employees  
WHERE department_id IS NOT NULL  
GROUP BY department_id  
HAVING COUNT(*) > 5;
```

As shown, removing records with no department_id reduce the data size



#7

Use Temporary Tables Wisely

WHAT IT IS ?

Strategic use of temporary tables for complex queries.

WHY IT MATTERS ?

Can simplify queries and improve performance in multi-step processes.

HOW TO OPTIMIZE ?


Use temporary tables to store intermediate results, especially when dealing with multiple complex joins or subqueries.



LET'S WORK WITH AN EXAMPLE

Before Optimization:


sql

 Copy code

```
SELECT e.first_name, e.last_name, d.department_name
FROM employees e
JOIN departments d ON e.department_id = d.department_id
JOIN locations l ON d.location_id = l.location_id
WHERE l.country_id = 'US';
```

After Optimization:

sql

 Copy code

```
-- Creating a temporary table for US departments
CREATE TEMP TABLE us_departments AS
SELECT d.department_id, d.department_name
FROM departments d
JOIN locations l ON d.location_id = l.location_id
WHERE l.country_id = 'US';

-- Joining employees with the temporary table
SELECT e.first_name, e.last_name, ud.department_name
FROM employees e
JOIN us_departments ud ON e.department_id = ud.department_id;
```



#8

Avoid or Optimize OR Clauses

WHAT IT IS ?

Replacing or optimizing OR clauses for efficiency.

WHY IT MATTERS ?

OR clauses can slow down queries; alternatives can improve performance.

HOW TO OPTIMIZE ?

Replace OR clauses with IN statements where possible, or break the query into multiple UNIONed queries.



LET'S WORK WITH AN EXAMPLE

Before Optimization:

sql

```
SELECT * FROM employees  
WHERE last_name = 'Smith' OR last_name = 'Jones';
```

After Optimization:

sql

```
SELECT * FROM employees  
WHERE last_name IN ('Smith', 'Jones');
```

As shown, replacing the OR Clause with IN would impact query speed.



Varun Sagar
Theegala

NEXT ➡

#9

Use EXISTS Instead of IN

WHAT IT IS ?

Using EXISTS for subquery checks instead of IN.

WHY IT MATTERS ?

EXISTS can be faster, especially with large subquery results.

HOW TO OPTIMIZE ?

Replace IN clauses with EXISTS when checking for the existence of a row in a subquery.



LET'S WORK WITH AN EXAMPLE

Before Optimization:

sql

```
SELECT e.first_name, e.last_name
FROM employees e
WHERE e.department_id IN
      (SELECT department_id
       FROM departments
       WHERE manager_id IS NOT NULL);
```

After Optimization:

sql

```
SELECT e.first_name, e.last_name
FROM employees e
WHERE EXISTS
      (SELECT 1
       FROM departments d
       WHERE e.department_id = d.department_id
       AND d.manager_id IS NOT NULL);
```



#10

Optimize Join Orders

WHAT IT IS ?

Ordering joins to process smaller datasets first.

WHY IT MATTERS ?

Reduces overall query processing time by minimizing intermediate result size.

HOW TO OPTIMIZE ?

Analyze the size of the tables involved in joins and structure the query to start with the smallest table, gradually joining larger tables.



LET'S WORK WITH AN EXAMPLE

Before Optimization:

sql

```
SELECT * FROM employees e
JOIN departments d ON e.department_id = d.department_id
JOIN locations l ON d.location_id = l.location_id;
```

After Optimization:

sql

```
SELECT * FROM locations l
JOIN departments d ON l.location_id = d.location_id
JOIN employees e ON d.department_id = e.department_id;
```

By replacing employees with a smaller dataset like locations as the base, we reduce intermediate table sizes



REMEMBER

Mastering these SQL optimization techniques is essential for data professionals aiming to **elevate data handling efficiency.**


Implementing these practices ensures your organization's data workflows are both **robust and responsive.**



Varun Sagar
Theegala

SHARE THIS

If you think
your network
would
find this
valuable



FOLLOW ME

I help you
**GROW &
SUSTAIN** as a
Data Analyst

